

Package ‘eimodel’

February 15, 2025

Title A fast alternative for the R x C ecological inference problem.

Version 0.0.0.1

Author Charles Thraves, Pablo Ubilla, Daniel Hermosilla

Maintainer Daniel Hermosilla <daniel.hermosilla.r@ug.uchile.cl>

Description Calculates the probability matrix of a R x C ecological inference problem with an Expected Maximization algorithm

License MIT + file LICENSE

Encoding UTF-8

Imports Rcpp,
R6,
jsonlite

LinkingTo Rcpp

NeedsCompilation yes

BugReports

<https://github.com/DanielHermosilla/ecological-inference-elections/issues>

Contents

eim	1
random_samples	12
simulate_elections	14
sum.eim	15
update.eim	16
Index	17

eim

An R6 Class for estimating an probability matrix from an R x C ecological inference problem.

Description

This class implements an EM algorithm using different methods to approximate the E-step, such as "Multinomial", "Hit and Run", "MVN CDF", "MVN PDF", and "Exact".

Usage

```
## S3 method for class 'eim'
summary(object)

## S3 method for class 'eim'
predict(object, ...)

## S3 method for class 'eim'
as.matrix(object)

std.eim(object, ...)
```

Public fields

<code>X (matrix)</code>	An (b x c) matrix with the observed results of the candidate votes (c) for each ballot box (b). Provided manually or loaded from JSON.
<code>W (matrix)</code>	An (b x g) matrix with the observed results of the demographic group votes (g) for each ballot box (b). Provided manually or loaded from JSON.
<code>method (character)</code>	A string indicating the EM method. One of: "Multinomial", "Hit and Run", "MVN CDF", "MVN PDF", or "Exact".
<code>probability (matrix)</code>	An (g x c) matrix that will store the final estimated probabilities of a given group (g) voting for a candidate (c).
<code>logLikelihood (numeric)</code>	A numeric vector that will store the log-likelihood values over all iterations of the Expectation-Maximization algorithm.
<code>total_iterations (integer(1))</code>	An integer indicating the total number of iterations performed by the Expectation-Maximization algorithm.
<code>total_time (numeric(1))</code>	The time taken by the EM algorithm.
<code>finish_state (character)</code>	The reason for the algorithm's termination. It can be either "Maximum iterations", "Log-likelihood decrease", "Convergence", or "Early exit".

Active bindings

<code>samples (integer(1))</code>	Active variable to show the Hit and Run samples if and only if the method is "Hit and Run".
<code>step_size (integer(1))</code>	Active variable to show the Hit and Run step size if and only if the method is "Hit and Run".
<code>multivariate_method (character)</code>	Active variable to show the method used to estimate the Multivariate Normal CDF, if and only if self\$method is "MVN CDF".

<code>multivariate_error (numeric(1))</code>	Active variable to show the error threshold for the Monte Carlo simulation of the Multivariate Normal CDF.
<code>multivariate_iterations (numeric(1))</code>	Active variable to show the number of iterations for the Monte Carlo simulation of the Multivariate Normal CDF.
<code>std (matrix)</code>	Active variable to show an estimate of the probability standard deviation if the bootstrapping method has been called.

Methods

Public methods:

- `eim$new()`
- `eim$precompute()`
- `eim$compute()`
- `eim$bootstrap()`
- `eim$print()`
- `eim$summary()`
- `eim$save_results()`
- `eim$clone()`

Method `new()`: Creates the object by defining the X and W matrix attributes.

Usage:

```
eim$new(X = NULL, W = NULL, json_path = NULL)
```

Arguments:

<code>X (matrix)</code>	A matrix (c x b) of observed candidate votes per ballot box (b) (optional; required if <code>json_path</code> is NULL).
<code>W (matrix)</code>	A matrix (b x g) of observed demographic group votes per ballot box (b) (optional; required if <code>json_path</code> is NULL).
<code>json_path (character)</code>	A string containing a path to a JSON file with "X" and "W" matrices. (optional; required if X or W are NULL)

Returns: An initialized eim object.

Examples:

```
# Example 1: Create a eim object from a matrix
model <- eim$new(X = matrix(1:9, 3, 3), W = matrix(1:9, 3, 3))

# Example 2: Create a eim object from a JSON file
\dontrun{
model2 <- eim$new(json_path = "a/file/path/to/a/file.json")
}
```

Method `precompute()`: Calculates the EM-independent variables for the Hit and Run and Exact methods.

Usage:

```
eim$precompute(method, ...)
```

Arguments:

<code>method</code> (<i>character</i>)	The method for precomputing. Options: "Hit and Run" or "Exact".
<code>...</code>	Additional arguments required by specific methods:

"Hit and Run" Method:

<code>step_size</code> (<i>integer(1)</i>)	The step size (M) for the Hit and Run algorithm. Must be a positive integer.
<code>samples</code> (<i>integer(1)</i>)	The number of samples (S) to generate. Must be an integer.

Returns: The modified eim object (for method chaining). Updates are made on the C internal memory.

Examples:

```
# Example 1: Precompute the Hit and Run method
simulations <- simulate_elections(num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20))
```

```
model <- eim$new(simulations$X, simulations$W)
```

```
model$precompute("Hit and Run",
  step_size = 1000,
  samples = 5000)
```

```
# Changes are made to the C internals API
```

```
# Example 2: Precompute the Exact method
```

```
model$precompute("Exact")
```

Method `compute()`: Executes the Expectation-Maximization (EM) algorithm based on the selected method. Additional parameters may be required depending on the method.

Usage:

```
eim$compute(
  method = "Multinomial",
  probability_method = "Group proportional",
  iterations = 1000,
  maximum_time = 1440,
  stopping_threshold = 0.001,
  verbose = FALSE,
  ...
)
```

Arguments:

<code>method (character)</code>	The method for estimating the Expectation-Maximization (EM) algorithm. Options: "Multinomial", "MVN CDF", "MVN PDF", "Hit and Run", and "Exact" (default: "Multinomial").
<code>probability_method (character)</code>	The method for obtaining the initial probability. Options: "Group proportional", "Proportional", or "Uniform". (default: "Group proportional").
<code>iterations (integer(1))</code>	The maximum number of iterations to perform on the EM algorithm. (default: 1000).
<code>maximum_time (integer(1))</code>	The maximum time (in minutes) that the algorithm will iterate. (default: 1440).
<code>stopping_threshold (numeric(1))</code>	The minimum difference between successive probabilities for stopping the iterations. (default: 0.001).
<code>verbose (boolean(1))</code>	A boolean indicating whether to print informative messages while iterating. (default: FALSE).
<code>...</code>	Additional arguments required by specific methods:
"Hit and Run" Method:	
<code>step_size (integer(1)):</code>	The step size (M) for the Hit and Run algorithm.
<code>samples (integer(1)):</code>	The number of samples (S) to generate.
"MVN CDF" Method:	
<code>multivariate_method (character):</code>	The integration method. Can be chosen between "Genz" and "Genz2". Default is "Genz2".
<code>multivariate_error (numeric(1)):</code>	The integration error threshold. Default is 1e-6.
<code>multivariate_iterations (integer(1)):</code>	The number of Monte Carlo iterations. Default is 5000.

Returns: The modified eim object (for method chaining).

Examples:

Example 1: Compute the Expectation-Maximization with default values

```

simulations <- simulate_elections(num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20))

model <- eim$new(simulations$X, simulations$W)
model$compute() # Retrieves the object with updated attributes

# Example 2: Compute the Expectation-Maximization for the Hit and Run method
\donttest{
  model$compute(method = "Hit and Run",
    step_size = 3000,
    samples = 1000)
}

# Example 3: Omit arguments to the Hit and Run method
\dontrun{
  model$compute(method = "Hit and Run",
    step_size = 3000)
  # Error; must supply the samples parameter as well
}

# Example 4: Run the MVN CDF with default values
model$compute(method = "MVN CDF")

# Example 5: Run an Exact estimation with user-defined parameters
model$compute(method = "Exact",
  probability_method = "Uniform",
  iterations = 5,
  stopping_threshold = 1e-3)
# Verbose was omitted

```

Method `bootstrap()`: Runs a bootstrap algorithm to estimate the **standard deviation** of probabilities.

Usage:

```

eim$bootstrap(
  bootstrap_iterations,
  ballot_boxes,
  method = "Multinomial",
  probability_method = "Group proportional",
  iterations = 1000,
  stopping_threshold = 0.001,
  verbose = FALSE,
  ...
)

```

Arguments:

<code>bootstrap_iterations</code> (<i>integer(1)</i>)	The number of EM computations.
<code>ballot_boxes</code> (<i>integer(1)</i>)	The number of ballot boxes to use as a sample. It must be

strictly less than the total
number of ballot boxes.

*Arguments for calling the `$compute()`
method:*

<code>method (character)</code>	The method for estimating the Expectation-Maximization (EM) algorithm. Options: "Multinomial", "MVN CDF", "MVN PDF", "Hit and Run", and "Exact" (default: "Multinomial").
<code>probability_method (character)</code>	The method for obtaining the initial probability. Options: "Group proportional", "Proportional", or "Uniform". (default: "Group proportional").
<code>iterations (integer(1))</code>	The maximum number of iterations to perform on the EM algorithm. (default: 1000).
<code>maximum_time (integer(1))</code>	The maximum time (in minutes) that the algorithm will iterate. (default: 1440).
<code>stopping_threshold (numeric(1))</code>	The minimum difference between successive probabilities for stopping the iterations. (default: 0.001).
<code>verbose (boolean(1))</code>	A boolean indicating whether to print informative messages while iterating. (default: FALSE).
...	Additional arguments required by specific methods:

"Hit and Run" Method:

<code>step_size (integer(1)):</code>	The step size (M) for the Hit and Run algorithm.
<code>samples (integer(1)):</code>	The number of samples (S) to generate.

"MVN CDF" Method:

<code>multivariate_method (character):</code>	The integration method. Can be chosen between "Genz" and "Genz2". Default is "Genz2".
<code>multivariate_error (numeric(1)):</code>	The integration error threshold. Default is 1e-6.

`multivariate_iterations (integer(1))`: The number of Monte Carlo iterations.
Default is 5000.

Returns: The modified eim object, under the field `$std`.

Examples:

```
simulations <- simulate_elections(num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20))
```

```
model <- eim$new(simulations$X, simulations$W)
```

```
model$bootstrap(30, 10)
```

```
model$std # An estimate of the standard deviation of the probabilities.
```

Method `print()`: Depending on the state of the algorithm (computed or not), it prints a message with its most relevant parameters.

Usage:

```
eim$print()
```

Returns: The object itself (for method chaining).

Examples:

```
simulations <- simulate_elections(num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20))
```

```
model <- eim$new(simulations$X, simulations$W)
print(model) # Prints the X and W matrices.
```

```
model$compute()
print(model) # Prints the X and W matrices along with the EM results.
```

Method `summary()`: Shows, in the form of a list, a selection of the most important attributes. It retrieves the method, number of candidates, ballots, and groups, as well as the principal results of the EM algorithm.

Usage:

```
eim$summary()
```

Returns: (*list*) A list with the method, candidates, ballots, groups, probabilities, and log-likelihood.

Examples:

```
simulations <- simulate_elections(num_ballots = 5,
  num_candidates = 3,
  num_groups = 2,
  ballot_voters = rep(100, 5))
```

```
model <- eim$new(simulations$X, simulations$W)
a_list <- model$summary()
a_list$method # Not computed yet
```



```
a_list$groups # 2
a_list$ballots # 5
names(a_list)
# "candidates" "groups" "ballots" "method" "probabilities" "logLikelihood"
```

Method `save_results()`: Saves the current eim object to a specified file. The results can be saved in:

- **RDS (Binary format)**: Preserves object structure for future use in R.
- **JSON**: Saves model data in a human-readable format.
- **CSV**: Saves the probability matrix in a tabular format.

Usage:

```
eim$save_results(filename)
```

Arguments:

<code>filename</code> (<i>character</i>)	The file name where the results should be saved, including its extension. The file extension determines the format:
<code>*.rds</code>	Saves as ‘RDS’ (default, binary format).
<code>*.json</code>	Saves as ‘JSON’ (readable and shareable).
<code>*.csv</code>	Saves the <i>final probability matrix</i> as ‘CSV’.

Returns: The modified eim object (for method chaining).

Examples:

```
simulations <- simulate_elections(num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20))

model <- eim$new(simulations$X, simulations$W)
model$compute()
model$save_results("results.rds") # Save as RDS
model$save_results("results.json") # Save as JSON
model$save_results("results.csv") # Save the final probability matrix as CSV
```

Note

Precomputing can eventually accelerate the computation time of the Expectation-Maximization algorithm.

The `$compute`, `$bootstrap`, `$print` and `$summary` methods can also be called as an S3 method with `predict()`, `std()`, `print()` and `summary()`.

Author(s)

Charles Thraves, Pablo Ubilla

References

Thraves, C. and Ubilla, P.: 'Fast Ecological Inference Algorithm for the RxC Case'

Examples

```
## -----
## Method `eim$new`
## -----

# Example 1: Create a eim object from a matrix
model <- eim$new(X = matrix(1:9, 3, 3), W = matrix(1:9, 3, 3))

# Example 2: Create a eim object from a JSON file
## Not run:
model2 <- eim$new(json_path = "a/file/path/to/a/file.json")

## End(Not run)

## -----
## Method `eim$precompute`
## -----

# Example 1: Precompute the Hit and Run method
simulations <- simulate_elections(num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20))

model <- eim$new(simulations$X, simulations$W)

model$precompute("Hit and Run",
  step_size = 1000,
  samples = 5000)
# Changes are made to the C internals API

# Example 2: Precompute the Exact method
model$precompute("Exact")

## -----
## Method `eim$compute`
## -----

# Example 1: Compute the Expectation-Maximization with default values

simulations <- simulate_elections(num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20))

model <- eim$new(simulations$X, simulations$W)
model$compute() # Retrieves the object with updated attributes

# Example 2: Compute the Expectation-Maximization for the Hit and Run method

model$compute(method = "Hit and Run",
  step_size = 3000,
```

```

    samples = 1000)

# Example 3: Omit arguments to the Hit and Run method
## Not run:
  model$compute(method = "Hit and Run",
    step_size = 3000)
# Error; must supply the samples parameter as well

## End(Not run)

# Example 4: Run the MVN CDF with default values
model$compute(method = "MVN CDF")

# Example 5: Run an Exact estimation with user-defined parameters
model$compute(method = "Exact",
  probability_method = "Uniform",
  iterations = 5,
  stopping_threshold = 1e-3)
# Verbose was omitted

## -----
## Method `eim$bootstrap`
## -----

simulations <- simulate_elections(num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20))

model <- eim$new(X = simulations$X, W = simulations$W)

model$bootstrap(30, 10)

model$std # An estimate of the probabilities' standard deviation.

## -----
## Method `eim$print`
## -----

simulations <- simulate_elections(num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20))

model <- eim$new(simulations$X, simulations$W)
print(model) # Prints the X and W matrices.

model$compute()
print(model) # Prints the X and W matrices along with the EM results.

## -----
## Method `eim$summary`
## -----

simulations <- simulate_elections(num_ballots = 5,
  num_candidates = 3,

```

```

    num_groups = 2,
    ballot_voters = rep(100, 5))

model <- eim$new(simulations$X, simulations$W)
a_list <- model$summary()
a_list$method # Not computed yet
a_list$groups # 2
a_list$ballots # 5
names(a_list)
# "candidates" "groups" "ballots" "method" "probabilities" "logLikelihood"

## -----
## Method `eim$save_results`
## -----

simulations <- simulate_elections(num_ballots = 20,
    num_candidates = 5,
    num_groups = 3,
    ballot_voters = rep(100, 20))

model <- eim$new(simulations$X, simulations$W)
model$compute()
model$save_results("results.rds") # Save as RDS
model$save_results("results.json") # Save as JSON
model$save_results("results.csv") # Save the final probability matrix as CSV

```

random_samples

Randomly create a voting instance by defining an interval

Description

Given a range of possible **observed** outcomes (such as the number of ballot boxes, candidates, etc.), it creates a completely random voting instance, simulating the unobserved results as well.

Usage

```

random_samples(
  ballots_range,
  candidates_range,
  group_range,
  voting_range,
  seed = NULL
)

```

Arguments

ballots_range (integer) A vector of length 2 specifying the lower and upper bounds for the number of ballot boxes.

candidates_range (integer) A vector of length 2 specifying the lower and upper bounds for the number of candidates to draw.

group_range (integer) A vector of length 2 specifying the lower and upper bounds for the number of demographic groups to draw.

voting_range	(integer) A vector of length 2 specifying the lower and upper bounds for the number of votes per ballot box.
seed	(<i>numeric(1)</i>) Optional. If provided, it overrides the current global seed.

Value

A list with components:

X	A matrix (b x c) containing candidate votes per ballot box.
W	A matrix (b x g) containing demographic votes per ballot box.
real_p	A matrix (g x c) containing the estimated (unobserved) probabilities that a demographic group votes for a given candidate.
ballots	The number of ballot boxes that were drawn.
candidates	The number of candidates that were drawn.
groups	The number of demographic groups that were drawn.
total_votes	A vector containing the total number of votes per ballot box.

See Also

[simulate_elections()]

Examples

```
bal_range <- c(30, 50)
can_range <- c(2, 4)
group_range <- c(2, 6)
voting_range <- c(50, 100)
results <- random_samples(bal_range, can_range, group_range, voting_range)

# X matrix
results$X # A randomly generated matrix of dimension (b x c)
ncol(results$X <= can_range[2]) # Always TRUE
ncol(results$X >= can_range[1]) # Always TRUE
nrow(results$X <= bal_range[2]) # Always TRUE
nrow(results$X >= bal_range[1]) # Always TRUE

# W matrix
results$W # A randomly generated matrix of dimension (b x g)
ncol(results$W <= group_range[2]) # Always TRUE
ncol(results$W >= group_range[1]) # Always TRUE
nrow(results$W <= bal_range[2]) # Always TRUE
nrow(results$W >= bal_range[1]) # Always TRUE

# Probability matrix
results$real_p # A matrix (g x c) that summarizes the unobserved outcomes
ncol(results$real_p) == ncol(results$X) # Always TRUE
nrow(results$real_p) == ncol(results$W) # Always TRUE
```

simulate_elections	<i>Simulates an election by creating the candidate and group matrices and their results.</i>
--------------------	--

Description

Given the number of ballots, groups, candidates, and votes per ballot, this function simulates an election. Additionally, it generates a resulting matrix (of dimension $g \times c$) representing the unobserved probabilities that a demographic group votes for a given candidate. These probabilities are drawn from a Dirichlet distribution and adjusted by a `lambda` value, which represents the heterogeneity of the groups.

Usage

```
simulate_elections(
  num_ballots,
  num_candidates,
  num_groups,
  ballot_voters,
  lambda = 0.5,
  seed = NULL
)
```

Arguments

<code>num_ballots</code>	<i>(integer(1))</i> The number of ballot boxes (" <i>b</i> ").
<code>num_candidates</code>	<i>(integer(1))</i> The number of candidates (" <i>c</i> ").
<code>num_groups</code>	<i>(integer(1))</i> The number of demographic groups (" <i>g</i> ").
<code>ballot_voters</code>	<i>(integer(num_ballots))</i> A vector of length <code>num_ballots</code> with the number of votes per ballot box.
<code>lambda</code>	<i>(numeric(1))</i> A value between 0 and 1 representing the heterogeneity of the groups. Values near 0 yield more heterogeneous results, but may be less realistic; the opposite is true for values closer to 1. (default: 0.5)
<code>seed</code>	<i>(numeric(1))</i> Optional. If provided, it overrides the current global seed. (default: NULL)

Value

A list with components:

<code>X</code>	A matrix ($b \times c$) with candidate votes per ballot box.
<code>W</code>	A matrix ($b \times g$) with demographic votes per ballot box.
<code>real_p</code>	A matrix ($g \times c$) with the estimated (unobserved) probabilities that a demographic group votes for a given candidate.

Note

Note: For the Dirichlet distribution, an `alpha` value of 1 yields a distribution that is uniform around the mean. Conversely, `alpha` values less than 1 tend to produce a sparser probability vector, with many values near zero and one or a few larger values. When `alpha` is greater than 1, the distribution becomes more concentrated around the mean.

See Also

[random_samples()]

Examples

```
result <- simulate_elections(num_ballots = 10,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = c(100, 10))
result$X # Candidate matrix (b x c)
result$W # Group matrix (b x g)
result$real_p # Probability matrix (g x c)
```

sum.eim	<i>Returns the total number of voters in the system.</i>
---------	--

Description

Given an initialized object, it returns the total number of voters. It is equivalent to running `sum(object$W)` or `sum(object$X)`.

Usage

```
## S3 method for class 'eim'
sum(object)
```

Arguments

object An eim object.

Value

The total number of voters.

Examples

```
simulations <- simulate_elections(
  num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20)
)

model <- eim$new(simulations$X, simulations$W)
sum(model) # 2000
```

update.eim

Update an existing eim model with a new EM algorithm computation

Description

This function updates an object by performing a new Expectation-Maximization computation with different parameters.

Usage

```
## S3 method for class 'eim'
update(object, ...)
```

Arguments

... New parameters to be passed to compute().
object An eim object.

Value

The updated object with the changes applied.

Examples

```
simulations <- simulate_elections(
  num_ballots = 20,
  num_candidates = 5,
  num_groups = 3,
  ballot_voters = rep(100, 20)
)

model <- eim$new(simulations$X, simulations$W)

predict(model)

update(model, "MVN PDF")

model$method # Returns 'MVN PDF'
```


Index

`as.matrix.eim(eim)`, [1](#)

`eim`, [1](#)

`predict.eim(eim)`, [1](#)

`random_samples`, [12](#)

`simulate_elections`, [14](#)

`std.eim(eim)`, [1](#)

`sum.eim`, [15](#)

`summary.eim(eim)`, [1](#)

`update.eim`, [16](#)